

Modeling Stock Market Returns using Deep Learning

[Midterm Report]

Devansh Gupta*

Computer Science and
Artificial Intelligence

IIIT Delhi

devansh19160@iiitd.ac.in

Divyansh Rastogi*

Computer Science and
Artificial Intelligence

IIIT Delhi

divyansh19464@iiitd.ac.in

Rupanshu Yadav*

Computer Science and
Artificial Intelligence

IIIT Delhi

rupanshu19475@iiitd.ac.in

1 Introduction

The stock market provides investors with the opportunity to share the profits of companies. More than five trillion dollars worth of stocks are traded every single day¹. Stock market forecasting has long been a popular and lucrative field of study, with more than 500 publications in the past five years (Hu et al., 2021a). Investors can maximize their profits by investing rationally based on technical analysis (de Souza et al., 2018). The time series prediction problem is complex due to the widely accepted semi-strong efficient market hypothesis (Malkiel and Fama, 1970) and their volatile & non-stationary structure (Adam et al., 2016). Hence features such as news sentiment, related markets & 8-K Reports are widely used. Prediction of assets or commodities such as currencies, gold, & crypto-currencies fall into the same categories of problems.

Stock market analysis methods fall into two categories, technical analysis, and fundamental analysis. Technical analysis involves the prediction of stock prices using historical data, primarily price and volume. It is worth noting that the more efficient the market is, the harder it would be to earn profits using technical analysis. On the other hand, the fundamental analysis makes predictions based on the intrinsic value of the investment by studying macroeconomic features, industry features, and company's features, including financial reports, competitors, & viability.

Statistical models such as Auto Regressive Moving Average (ARMA), Auto-Regressive Integrated Moving Average (ARIMA), Smooth Transition Autoregressive (STAR) are widely used for time series prediction purely based on the historical price of a stock. Machine learning has been intensively

researched for its use in financial market forecasting. Decision trees, XGBoost, and Naïve Bayes all have shown to be beneficial in stock market price prediction (Ballings et al., 2015). Deep learning methods have also proven to be helpful in time series analysis (Fawaz et al., 2019).

2 Related Work

Financial markets have shown to be volatile and non-stationary (Dionisio et al., 2007; Adam et al., 2016) with having a close resemblance to an ensemble of particles in statistical mechanics, precisely Brownian motion (Osborne, 1959). Despite the existence of multiple opposing views on the efficiency of stock markets, research reveals that financial markets are, to some extent, predictable (Ferreira and Santa-Clara, 2011; Kim et al., 2011; Bollerslev et al., 2014). Studies have evaluated the predictive relationships of numerous financial and economic variables (Enke and Thawornwong, 2005; Kim and Han, 2000; Vellido et al., 1999), showing that the closer historical data time is to the present, the stronger the data's impact on the predictive model (Liao and Wang, 2010).

Various statistical & econometric machine learning methods (Agrawal et al., 2013) have been employed to predict stock returns/directions like ARIMA, SVM (Schumaker and Chen, 2009; Lee, 2009; Cao and Tay, 2001), Boosted Regression Trees (Pierdzioch et al., 2018) and Logistic Regression (Chong et al., 2017). Ou and Wang 2009 discussed & applied ten different data mining techniques to predict price movement in the Hong Kong stock market Hang Seng index. For a survey on various employed machine learning techniques for stock market predictions, we refer the reader to Strader et al. 2020.

With the rise of deep learning applications in stock markets (Hu et al., 2021b), its shown to

¹<https://www.nasdaq.com/articles/forex-market-overview-2019-06-07>

help model complex intrinsic relationships and extract abstract features from data without relying on econometric assumptions or human expertise (Chong et al., 2017). Although to achieve accurate results with neural networks, it is important to have a deliberate selection of input variables (Lam, 2004). For market prediction, various feature selection & extraction methods have been used alongside machine learning techniques such as Genetic Algorithms (Kim and Han, 2000), Auto-encoders (Chong et al., 2017), Restricted Boltzmann Machine (RBM) (Chong et al., 2017) and Principal Component Analysis (PCA) (Zhong and Enke, 2019; Chong et al., 2017). Zhong and Enke 2019 provides a comprehensive study on PCA and its non-linear fuzzy variants used along with Artificial Neural Networks (ANN) and Deep Neural Networks (DNN) for forecasting the daily return direction of the SPDR S&P 500 ETF index.

3 Dataset

We used the CNNPred dataset used in (Hoseinzade and Haratizadeh, 2019) to get results on the existing models. This dataset consisted of market prices of NASDAQ, NYSE, S&P500, DJI, and RUSSELL from 2010 to 2017. This dataset had the market's closing price at a particular day along with a diverse set of features consisting of technical indicators like Exponentially Weighted Average of data from a varying number of previous days, returns, and their weighted averages. Along with the technical indicators, the dataset consisted of additional data on commodities like oil, gold, and silver prices. The features also had the currency exchange rates, returns from other markets, and stock futures which have shown to be important in determining the stock's current price on the market.

4 Baselines

We have chosen our baselines to explore a wide range of machine learning methods to stock price prediction and trend classification. Our baselines start with statistical models, transitioning into classical machine learning, and then to deep learning.

4.1 ARIMA

ARIMA is widely used for prediction in time series data modeling. ARMA requires the time series to be stationary; which means a constant mean, constant variance and non-seasonal. Since, stock market prices don't have constant mean, hence a

transformed feature, defined as

$$y_i = price_i - price_{i-d}$$

where d is the lag. Auto-Regressive(AR) and Moving Average(MA) models are parameterized by auto-correlation coefficient(p) and Partial auto-correlation coefficient(q). A search is performed given the maximum values of p and q to find the best model. An RMSE of 423.24 was observed. However, if the first testing point is included in the training data to predict the next point, an RMSE of 49.72 was observed.

4.2 PCA+ANN

This model has been given in Zhong and Enke 2019 for forecasting daily return direction of the SPDR S&P 500 ETF index. Upon examination, we discovered that all of the 60 financial variables evaluated in the study's dataset are already present in & analogous to the 82 features of our chosen dataset, CNNPred. Thus, the study's model is extrapolated to the CNNPred dataset's S&P market.

Zhong and Enke explored multiple data transformation techniques including PCA and its variants, fuzzy robust principal component analysis (FR-PCA) and kernel-based principal component analysis (KPCA), among others. Their results showed that traditional PCA outperformed all non-linear techniques on real-world data. Thus, PCA is chosen as our data transformation technique and PCA-represented dataset with 82 principal components is used.

4.2.1 Data Preprocessing

A classical statistical principle is used for detection of outliers based on inter-quartile ranges (Navidi, 2011). These outliers are accordingly adjusted similar to a method used by Cao and Tay 2001. The cleaned data is split in 70/15/15 ratio for train, validation and test dataset respectively. The data is standardized with the mean and variance of the training dataset.

4.2.2 Model & Training

The PCA-represented dataset is classified using a ANN network comprising of 4 layers with RELU activation & Sigmoid activation for the last layer. Dropout has been introduced in the network to avoid overfitting. Binary cross entropy loss is used as the loss criterion. The initial learning rate was set to 0.0001 with ADAM optimizer for training

over maximum of 100 epochs. Early stopping is implemented with the use of validation set.

4.2.3 Results

We obtained an accuracy of 0.559 and a F1-score of 0.656 on the test set.

4.3 CNN-Pred2D

This model (Hoseinzade and Haratizadeh, 2019) classifies the change in the closing price of the market using only the data for the market under analysis. It takes the input of all the features from last 60 days and leverages 2D convolution filters for making feature maps and finally classifying the change in price. Since CNNs are good at capturing short range data and hierarchically extract features from day wise data, they serve as a good method for predicting stock prices. The initial learning rate for training this network was set to 0.001, and the ADAM optimizer was used for training over 100 epochs. Since the network was quite small with only one fully connected layer and three convolution layers, a weight decay of 0.0001 was sufficient to regularize the network. Each feature was normalized according to the training set and was used during the validation and test time.

Market Name	Average	Maximum
NASDAQ	0.438	0.552
NYSE	0.491	0.492
S&P500	0.434	0.573
RUSSELL	0.497	0.691
DJI	0.471	0.696

Table 1: F-Scores on the test set using CNNPred2D trained for 100 Epochs on the CNNPred Stock market Dataset

4.4 CNN-Pred3D

This model (Hoseinzade and Haratizadeh, 2019) classifies the change in the closing price of the market using the data of various markets. It takes the input as a 3D block of all the features from the 5 markets in the dataset over the last 60 days and leverages 3D convolution filters for making feature maps and finally classifying the change in price by using data across many markets. The usage of 3D convolution filters is the same as 2D convolution filters only that these filters hierarchically extract features and are temporally sensitive over an additional dimension. Similar to the CNN-Pred2D, the initial learning rate for training this network was set to 0.001, and the ADAM optimizer was used

for training over 100 epochs with a weight decay of 0.0001. Each feature was normalized according to the training set and was used during the validation and test time.

Market Name	Average	Maximum
NASDAQ	0.49	0.536
NYSE	0.432	0.566
S&P500	0.491	0.67
RUSSELL	0.489	0.521
DJI	0.486	0.5

Table 2: F-Scores on the test set using CNNPred3D trained for 100 Epochs on the CNNPred Stock market Dataset

4.5 CNN-LSTM Model

This model (Lu et al., 2020) considers learning sequence aware features as the stock market is an event which moves in the temporal dimension, thus it is difficult to ignore the sequential information present in the latent embeddings for the downstream tasks. This work explored the price prediction perspective for more informed stock market trading and hence was a regression task. The convolutional features were used to calculate the temporal features over time stamps and then an LSTM was used to capture the sequential features. The structure of LSTM is designed in such a manner that it works on selectively learning which information to hide and which to infer on over a certain time step and pass both the states for the next time step (Hochreiter and Schmidhuber, 1997). Thus, this overparameterization leads to a delayed stability of LSTMs in terms of metrics but provably results to a more optimal result in lesser number of iterations (Arora et al., 2018). This model gave sufficiently good results on predicting the closing price of a market on training with 100 epochs with a learning rate of 0.001 on ADAM optimizer, with a weight decay of 0.0001.

Market Name	MAE	RMSE	R^2
NASDAQ	645.21	854.89	0.99
NYSE	433.31	577.70	0.99
S&P500	143.58	192.79	0.99
DJI	162.90	205.95	0.99
RUSSELL	1899.07	2551.76	0.99

Table 3: Regression metrics on the test set using CNN-LSTM model trained for 100 Epochs on the CNNPred Stock market Dataset

5 Conclusions and Future Work

There has been work done that shows that the forward propagation of residual networks is a discrete solution for ODEs and proposes an efficient method to learn the dynamics of the system using ODE solving (Chen et al., 2018). We plan to explore this direction of work as it has shown to model irregular time series data very well and may capture the dynamics lightly corresponding to Brownian Motion in the stock market (Osborne, 1959).

Generative Adversarial Networks can map the distribution of data to a prior. We plan on further exploring the concepts of GAN and link it with financial data (Takahashi et al., 2019).

We also plan to explore the company relations as an additional form of supervision to use models like Graph Attention Networks or Hypergraph Convolutions based on the work done in (Sawhney et al., 2021). We further plan to explore the use of ARIMA as input features while exploring the above deep learning models.

6 Contributions

All authors denoted by * contributed equally, the ordering of naming is alphabetical.

References

- Klaus Adam, Albert Marcet, and Juan Pablo Nicolini. 2016. Stock Market Volatility and Learning. *The Journal of Finance*, 71(1):33–82.
- Janhavi Agrawal, Vijay S. Chourasia, and A. K. Mitra. 2013. State-of-the-art in stock prediction techniques. *International Journal of Advanced Research in Electrical, Electronics and Instrumentation Energy*, 2:1360–1366.
- Sanjeev Arora, Nadav Cohen, and Elad Hazan. 2018. On the optimization of deep networks: Implicit acceleration by overparameterization. In *Proceedings of the 35th International Conference on Machine Learning*, volume 80 of *Proceedings of Machine Learning Research*, pages 244–253. PMLR.
- Michel Ballings, Dirk Van den Poel, Nathalie Hespeels, and Ruben Gryp. 2015. *Evaluating multiple classifiers for stock price direction prediction*. *Expert Systems with Applications*, 42(20):7046–7056.
- Tim Bollerslev, James Marrone, Lai Xu, and Hao Zhou. 2014. *Stock return predictability and variance risk premia: Statistical inference and international evidence*. *Journal of Financial and Quantitative Analysis*, 49(3):633–661.
- Lijuan Cao and Francis Tay. 2001. *Financial forecasting using support vector machines*. *Neural Computing and Applications*, 10:184–192.
- Ricky T. Q. Chen, Yulia Rubanova, Jesse Bettencourt, and David K Duvenaud. 2018. Neural ordinary differential equations. In *Advances in Neural Information Processing Systems*.
- Eunsuk Chong, Chulwoo Han, and Frank Park. 2017. *Deep learning networks for stock market analysis and prediction: Methodology, data representations, and case studies*. *Expert Systems with Applications*, 83.
- Andreia Dionisio, Rui Menezes, and Diana A. Mendes. 2007. On the integrated behaviour of non-stationary volatility in stock markets. *Physica A: Statistical Mechanics and its Applications*, 382(1):58–65. Applications of Physics in Financial Analysis.
- David Enke and Suraphan Thawornwong. 2005. *The use of data mining and neural networks for forecasting stock market returns*. *Expert Systems with Applications*, 29:927–940.
- Hassan Ismail Fawaz, Germain Forestier, Jonathan Weber, Lhassane Idoumghar, and Pierre-Alain Muller. 2019. *Deep learning for time series classification: a review*. *Data Mining and Knowledge Discovery*, 33(4):917–963. ArXiv: 1809.04356.
- Miguel A. Ferreira and Pedro Santa-Clara. 2011. *Forecasting stock market returns: The sum of the parts is more than the whole*. *Journal of Financial Economics*, 100(3):514–537.
- Sepp Hochreiter and Jürgen Schmidhuber. 1997. Long short-term memory. *Neural computation*, 9(8):1735–1780.
- Ehsan Hoseinzade and Saman Haratizadeh. 2019. Cnnpred: Cnn-based stock market prediction using a diverse set of variables. *Expert Systems with Applications*, 129:273–285.
- Zexin Hu, Yiqi Zhao, and Matloob Khushi. 2021a. *A Survey of Forex and Stock Price Prediction Using Deep Learning*. *Applied System Innovation*, 4(1):9.
- Zexin Hu, Yiqi Zhao, and Matloob Khushi. 2021b. *A survey of forex and stock price prediction using deep learning*. *Applied System Innovation*, 4(1).
- Jae H. Kim, Abul Shamsuddin, and Kian-Ping Lim. 2011. Stock return predictability and the adaptive markets hypothesis: Evidence from century-long u.s. data. *Journal of Empirical Finance*, 18(5):868–879.
- Kyoung-jae Kim and Ingoo Han. 2000. Genetic algorithms approach to feature discretization in artificial neural networks for the prediction of stock price index. *Expert Systems with Applications*, 19:125–132.

- Monica Lam. 2004. [Neural network techniques for financial performance prediction: integrating fundamental and technical analysis](#). *Decision Support Systems*, 37(4):567–581. Data mining for financial decision making.
- Ming-Chi Lee. 2009. [Using support vector machine with a hybrid feature selection method to the stock trend prediction](#). *Expert Systems with Applications*, 36:10896–10904.
- Zhe Liao and Jun Wang. 2010. [Forecasting model of global stock index by stochastic time effective neural network](#). *Expert Systems with Applications*, 37:834–841.
- Wenjie Lu, Jiazheng Li, Yifan Li, Aijun Sun, and Jingyang Wang. 2020. A cnn-lstm-based model to forecast stock prices. *Complexity*, 2020:6622927.
- Burton G. Malkiel and Eugene F. Fama. 1970. [Efficient Capital Markets: A Review of Theory and Empirical Work*](#). *The Journal of Finance*, 25(2):383–417.
- Navidi. 2011. Statistics for engineers and scientists. *McGraw-Hill, New York*, 3rd edn.
- M. F. M. Osborne. 1959. Brownian motion in the stock market. *Operations Research*, 7(2):145–173.
- Ou and Wang. 2009. [Prediction of stock market index movement by ten data mining techniques](#). *Modern Applied Science*, 3.
- Christian Pierdzioch, Rangan Gupta, Hossein Hassani, and Emmanuel Silva. 2018. [Forecasting Changes of Economic Inequality: A Boosting Approach](#). Working Papers 201868, University of Pretoria, Department of Economics.
- Ramit Sawhney, Shivam Agarwal, Arnav Wadhwa, Tyler Derr, and Rajiv Ratn Shah. 2021. Stock selection via spatiotemporal hypergraph attention network: A learning to rank approach. In *AAAI*.
- Rob Schumaker and Hsiu-chin Chen. 2009. [Textual analysis of stock market prediction using breaking financial news: The azfin text system](#). *ACM Trans. Inf. Syst.*, 27.
- Matheus José Silva de Souza, Danilo Guimarães Franco Ramos, Marina Garcia Pena, Vinicius Amorim Sobreiro, and Herbert Kimura. 2018. [Examination of the profitability of technical analysis based on moving average strategies in BRICS](#). *Financial Innovation*, 4(1):3.
- Troy J. Strader, John J. Rozycki, Thomas H. Root, and Yu-Hsiang Huang. 2020. Machine learning stock market prediction studies: Review and research directions. *Journal of International Technology and Information Management*, 28:63–83.
- Shuntaro Takahashi, Yu Chen, and Kumiko Tanaka-Ishii. 2019. Modeling financial time-series with generative adversarial networks. *Physica A: Statistical Mechanics and its Applications*, 527:121261.
- A. Vellido, P.J.G. Lisboa, and K. Meehan. 1999. [Segmentation of the on-line shopping market using neural networks](#). *Expert Systems with Applications*, 17(4):303–314.
- Xiao Zhong and David Enke. 2019. [Predicting the daily return direction of the stock market using hybrid machine learning algorithms](#). *Financial Innovation*, 5.